


Lecture 2: Data Representation I – Bits, Bytes & Integers

Session 2 · Mon Aug 31, 2026

From Physical Memory to One Byte

That is a **RAM module** — physical hardware, many chips, not one byte.

Memory is **byte-addressable**: every address names one **8-bit byte**, and a byte is the smallest unit a C program can address on its own.



One byte = 8 bits. Today is about what fits in one, and how to work on it.

One Byte, Two Readings

Start with one byte: eight switches that are each either 0 or 1.

One byte: the bit pattern for 65



64 + 1 = 65


as text (ASCII): A

The computer stores only the bits; the program decides how to read them.

Counting in Binary

Why does **01000001** mean 65? Because each place is worth twice the one to its right — the same idea as decimal's ones, tens and hundreds, with 2 instead of 10.

Each place is worth twice the place to its right



Turn a bit on and you add its place value.

$$128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = 255$$

Counting Up

Counting works exactly the way it does in decimal — when a column runs out of digits, it rolls over and carries into the next one:

0 = 0000	3 = 0011	6 = 0110	9 = 1001
1 = 0001	4 = 0100	7 = 0111	10 = 1010
2 = 0010	5 = 0101	8 = 1000	11 = 1011

Each extra bit doubles the patterns available: 8 bits give $2^8 = 256$. So an unsigned byte runs 0–255, and an N-bit unsigned value runs 0 to $2^N - 1$ — its **UMax**, modulo 2^N .

Writing the Same Number Three Ways

A number does not change when you write it differently. C lets you say which notation you are using with a **prefix**:

Written	Prefix means	Value
42	no prefix — plain decimal	forty-two
<code>0b101010</code>	<code>0b</code> — read the digits as binary	forty-two
<code>0x2A</code>	<code>0x</code> — read the digits as hexadecimal	forty-two

All three are the same value to the compiler. You pick whichever makes the thing you are doing readable: decimal for counting, binary for seeing bits, hex for writing bytes compactly.


Note

`0b` is not standard until C23. GCC and Clang accept it, and these slides use it to show bit patterns, but in ISO C17 write `0x2A` or `42`.

Why Hex Fits Bytes So Well

Eight ones and zeros is easy to miscount. Four bits count 0–15 — exactly one hex digit — so one byte is always two digits, and you can still read the bits straight off the page.

One byte = two hex digits



**4 bits count 0 to 15,
so we need 6 extra digits:**

10 = A
11 = B
12 = C
13 = D
14 = E
15 = F

0 through 9 stay as they are.

Characters Are Just Numbers

The rule that turns 65 into **A** is **ASCII**: a table agreed in the 1960s giving every character a number from 0 to 127.

Characters	Codes	Worth knowing
'0'-'9'	48-57 (0x30-0x39)	consecutive, so <code>c - '0'</code> gives the digit
'A'-'Z'	65-90 (0x41-0x5A)	consecutive
'a'-'z'	97-122 (0x61-0x7A)	exactly 32 more than the capitals

```
char c = '7';
int d = c - '0';           // 7
int is_digit = (c >= '0' && c <= '9'); // 1
```

In C the literal **'A'** is the number 65. Nothing clever is happening.

The Same Ideas in Python

C

```
int x = 65;
printf("%d\n", x);      // 65
printf("%c\n", x);      // A

unsigned a = 12, b = 10;
printf("%u\n", a & b);  // 8
printf("%u\n", a << 2); // 48
```

Python

```
x = 65
print(x)          # 65
print(chr(x))     # A

a, b = 12, 10
print(a & b)       # 8
print(a << 2)     # 48
```

The operators are the same. The difference is the box: a C value has a fixed width, and a Python integer just grows.

Bit-Level Operations

For these bit patterns, C provides six bitwise operators:

Op	Name	Example
<code>&</code>	AND	<code>0b1100 & 0b1010 = 0b1000</code>
<code>\ </code>	OR	<code>0b1100 \ 0b1010 = 0b1110</code>
<code>^</code>	XOR	<code>0b1100 ^ 0b1010 = 0b0110</code>
<code>~</code>	NOT	<code>~0b1100 = 0b1111...11110011</code> (every bit flips, including the leading zeros)
<code><<</code>	Left shift	<code>0b0011 << 2 = 0b1100</code>
<code>>></code>	Right shift	<code>0b1100 >> 2 = 0b0011</code>

Warning

Right shift of signed integers is **implementation-defined** in C. GCC uses **arithmetic** right shift (preserves sign bit). Don't rely on this unless you're sure of your compiler.

Bitwise Operations, One Column at a Time

The same two inputs, three separate rules

Every output bit is decided on its own, from just the two bits directly above it.

INPUTS

a	1	1	0	0	= 1100
b	1	0	1	0	= 1010

a & b	1	0	0	0	AND = 1000
-------	---	---	---	---	------------

1 only where both are 1

a b	1	1	1	0	OR = 1110
-------	---	---	---	---	-----------

1 where either is 1

a ^ b	0	1	1	0	XOR = 0110
-------	---	---	---	---	------------

1 where they differ

De Morgan's Laws

Two rules let you rewrite any mix of $\&$, $|$ and $\sim$ — and they are how Lab 1 asks you to build $\wedge$ out of nothing but $\sim$ and $\&$.

Flip the whole thing, and AND becomes OR

$$\sim(x \ \& \ y) == \sim x \ | \ \sim y$$

$$\sim(x \ | \ y) == \sim x \ \& \ \sim y$$


x	y	x & y	$\sim(x \ \& \ y)$	$\sim x \ \ \sim y$
0	0	0	1	1
0	1	0	1	1
1	0	0	1	1
1	1	1	0	0

same in all four rows

Two inputs give four possible rows, so checking all four is a complete proof.

Shifting Bits

Shifting slides every bit sideways




Bits pushed off the end are gone for good; zeros come in behind them.

Shifting left by n multiplies by 2^n . Shifting an unsigned value right by n divides by 2^n .

Test, Set, Clear, Toggle

`1u << 3` builds a mask with exactly one bit set



Bit 0 is the rightmost — the ones place. The mask is 0 everywhere except the bit you want.

```
unsigned x = 0xAE;           // 10101110
x |= (1u << 3);              // SET    – bit 3 becomes 1
x &= ~(1u << 3);             // CLEAR  – bit 3 becomes 0
x ^= (1u << 3);              // TOGGLE  – bit 3 flips
if (x & (1u << 3)) { }      // TEST   – nonzero when bit 3 is 1
```

Those four are most of Lab 1. Keep the value **unsigned** — shifting into a sign bit is undefined.

What Two's Complement Means

Complement means the part that completes a whole. Flip every bit of x to get its *one's complement*; add 1 to that and you have its **two's complement**, which is how C writes $-x$:

To write -5 in 8 bits

1. Start with +5

00000101

↓ flip every 0 ↔ 1

2. One's complement

11111010

+ 1

↓

3. Two's complement

11111011

= -5

More 8-bit examples

+1: 00000001

-1: 11111111

0: 00000000

two's: 00000000


Zero stays zero.

A ninth bit is simply discarded — 8 bits keep only their rightmost 8.

The Range It Buys You

An N-bit two's complement covers $-2^{(N-1)}$ to $2^{(N-1)} - 1$:

8-bit two's-complement values



There is one more negative value than positive value.

$TMin = -TMax - 1$ — there is no positive 128 in 8 bits, so negating $TMin$ overflows and is undefined.

Sign Extension

When converting a smaller signed type to a larger one, copy the sign bit:

```
int8_t  x = -1;    // 0b11111111  
int16_t y = x;     // 0b1111111111111111 = -1 ✓
```

Copy the sign bit into the new high bits

int8_t -1

11111111

extend



int16_t -1

1111111111111111

Truncation

When converting larger to smaller, drop the high bits:

```
int16_t  x = 257;  
uint8_t  y = x;    // 1: conversion to unsigned is modulo 256
```

Keep the low bits; discard the high bits

int16_t 257

000000001

000000001



keep these low 8 bits

int8_t result = 00000001 = 1

Size of Integer Types

```
printf("char:      %zu\n", sizeof(char));      // 1 C byte (not always 8 bits)
printf("short:     %zu\n", sizeof(short));     // 2
printf("int:       %zu\n", sizeof(int));       // 4
printf("long:      %zu\n", sizeof(long));      // 8 (LP64)
printf("long long: %zu\n", sizeof(long long)); // 8
printf("pointer:   %zu\n", sizeof(void*));     // 8
```




Tip

Use `sizeof` and fixed-width types (`int32_t`, `uint64_t`) when exact sizes matter. Never assume `int` is 4 bytes; use `CHAR_BIT` when bit width matters. Plain `char` is its own trap: whether it is signed is implementation-defined, so say `signed char` or `unsigned char` when it matters.


Byte Ordering (Endianness)

An `int` occupies 4 consecutive bytes, so those 4 bytes have to be laid out in some order. **Endianness** is that order. One question matters: **which byte goes at the lowest address?**

Little-endian: D is first in memory



Big-endian: A is first in memory



```
int x = 1;    // ask the machine: which byte of x sits at the lowest address?
if (*(unsigned char *)&x == 1) printf("Little-endian\n");
else                               printf("Big-endian\n");
```

AI Systems Connection

Quantization is one important reason large models can fit on more hardware. A common symmetric int8 scheme represents a block of weights with 8-bit signed values and a scale factor:

$$\text{real_value} \approx \text{scale} \times \text{int8_value} \qquad \text{scale} = \text{max_abs} / 127$$

Everything today shows up here. Two int8 values need 16 bits for their product, so a long dot product accumulates into something wider — often `int32_t`. And if a signed weight is loaded as *unsigned* into a wider register, sign extension does not happen, negatives become huge positives, and the answer is quietly wrong.

Try it: quantize a float array with `scale = max_abs / 127`, dequantize, and measure the error.

Practice

1. Write `0xD6` in binary, and `10110010` in hex, without using a computer.
2. What does `x & 1` tell you about `x`? And what does `x & (x - 1)` do? Try both on `0b0110` by hand.
3. Write one expression that clears bit 5 of `unsigned x`, and one that tests whether bit 0 and bit 7 are both set.
4. What is `TMax` for a 32-bit `int`? What is `TMin`?
5. Write a C expression that checks if your machine is little-endian.
6. Is left-shifting a 32-bit `int` by 32 positions defined in C? Why not?
7. Why is `-TMin` not representable in the same signed type?

Reading

- CS:APP §2.1 (Information Storage)
- CS:APP §2.2 (Integer Representations)